The Many-Path Mind

Rethinking How AI Agents Choose

ABSTRACT

"A single answer is not a solution — it is a guess that got lucky."

Modern AI agents are structurally fragile: commit to one execution path and there is no mechanism to recover when that path is wrong. This report examines **EnCompass** — a framework built on **Probabilistic Angelic Nondeterminism (PAN)** that separates program logic from decision strategy, allowing agents to explore multiple execution paths before selecting the best outcome. Tested across code translation, hypothesis generation, and complex reasoning, it delivers measurable gains in accuracy and reliability.

1. THE PROBLEM WITH SINGLE-PATH EXECUTION

Most AI agents operate like a student who writes the first answer that comes to mind and submits without review. One output, no comparison, no recovery mechanism. If the first attempt is wrong, the system has no structural way to find out.

This is more than inconvenience — it is architectural. Mixing what a program should do with how it decides conflates two separate concerns, making systems brittle, hard to debug, and difficult to improve. EnCompass proposes a clean separation, and that separation turns out to matter considerably.

> *"Between the first attempt and the best answer lies a space most AI systems never enter."*

2. PROBABILISTIC ANGELIC NONDETERMINISM

The conceptual engine behind EnCompass is PAN. Rather than forcing a program to commit to a single action sequence, PAN allows it to consider multiple possible ways of reaching the goal simultaneously — then selects the one that performs best.

The "angelic" framing assumes the system always picks the best available choice, giving it a principled way to explore a space of executions without becoming computationally intractable. This is a philosophical reframing as much as a technical one: uncertainty becomes something to navigate deliberately, not something to resolve immediately with a guess.

3. HOW ENCOMPASS WORKS

The framework is implemented in Python and designed to integrate naturally with large language models. Its four-step workflow is cleanly separated:

APPLICATIONS & RESULTS

Tested across three deliberately varied domains:

Code Translation

Java→Python translation. Multi-path execution found better structural equivalences than single-pass methods.

Hypothesis Generation

Multiple candidate ideas generated and compared; the most coherent was selected automatically.

Reasoning Tasks

Complex multi-step problems solved more reliably by comparing reasoning chains rather than committing to the first.

Improvements held across all three domains, suggesting the gains come from the architecture itself — not from domain-specific tuning.

ADVANTAGES

+ Accuracy improves across code, reasoning, and hypothesis tasks
+ Logic and decision-making cleanly separated — easier to maintain
+ Search strategy can be upgraded without touching task logic

LIMITATIONS

− Requires significantly more compute than single-path methods
− Poor search design can harm rather than help performance
− Implementation complexity is higher than traditional pipelines

01 User writes a normal program no special syntax required.

02 EnCompass transforms it into multiple possible execution paths.

03 A dedicated search layer evaluates each path on performance or correctness.

04 The best result is returned. Decision-making stays separate from task logic.

That modularity is significant: the search strategy can be improved independently without touching the program logic.

"The difference between a system that answers and a system that searches for answers is not speed — it is honesty about uncertainty."

REFERENCES

1. EnCompass: Enhancing Agent Programming with Search-Based Execution. Research Paper Summary (2026).
2. Macmillan-Scott, O., & Musolesi, M. (2025). (Ir)rationality in AI. arXiv:2311.17165.
3. Kaelbling, L., & Hedden, B. (2026). The philosophical puzzle of rational artificial intelligence. MIT News.